

Marc Vintró Alonso

# DevOps Test Notes

I installed the provided helm package [ping.tar.gz](#) in GCP to test the functionality.

I created a cluster with the following parameters **[1]**.

Cluster was created **[2]**.

I manually installed the helm chart using the cloud shell.

```
$tar -zxvf ping.tar.gz
```

I connected to the cluster to get the cluster credentials :

```
$gcloud container clusters get-credentials cluster-1 --zone us-central1-a --project instant-jetty-435813-p4
```

and installed the helm chart on it :

```
$helm install ping ./ --namespace default
```

I got the following error:

```
Error: INSTALLATION FAILED: parse error at (ping/templates/deployment.yaml:51): unexpected EOF
```

I added `{{- end }}` to `deployment.yaml` to be able to deploy it.

```
NAME: ping
```

```
LAST DEPLOYED: Sat Jan 18 10:32:33 2025
```

```
NAMESPACE: default
```

```
STATUS: deployed
```

```
REVISION: 1
```

```
NOTES:
```

```
1. Get the application URL by running these commands:
```

```
export POD_NAME=$(kubectl get pods --namespace default -l
"app.kubernetes.io/name=ping,app.kubernetes.io/instance=ping" -o jsonpath="{.items[0].metadata.name}")
export CONTAINER_PORT=$(kubectl get pod --namespace default $POD_NAME -o
jsonpath="{.spec.containers[0].ports[0].containerPort}")
echo "Visit http://127.0.0.1:8080 to use your application"
kubectl --namespace default port-forward $POD_NAME 8080:$CONTAINER_PORT
```

Pod is up and running :

```
$kubectl get pods
```

```
NAME                                READY STATUS RESTARTS AGE
ping-5989f667d5-624mc 1/1 Running 0 1m02s
```

After forwarding the port I was able to see the app running **[3]**.

# -----CHALLENGE 1-----

“Modify the Ping Helm Chart to deploy the application on the following restrictions:  
Isolate specific node groups forbidding the pods scheduling in these node groups.

Ensure that a pod will not be scheduled on a node that already has a pod of the same type. Ensure that Pods are deployed across different availability zones. Ensure that another random service is up before applying the manifests.”

Isolate Specific Node Groups :

I will prevent Prevent pods from scheduling on specific node groups by using node affinity[4]. In templates/deployment.yaml, I updated the spec.template.spec section to include affinity. Here's the modified section:

```
spec:
  {{- with .Values.nodeSelector }}
  nodeSelector:
    {{- toYaml . | nindent 8 }}
  {{- end }}
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: "node-group"
                operator: NotIn
                values:
                  {{- toYaml .Values.forbiddenNodeGroups | nindent 12 }}
```

I updated values.yaml, I added a new configuration field for forbiddenNodeGroups under values.yaml:

```
# List of node groups where pods are not allowed to be scheduled
forbiddenNodeGroups:
  - forbidden-group-1
  - forbidden-group-2
```

**nodeAffinity:** Specifies the scheduling rules for the pods.

**requiredDuringSchedulingIgnoredDuringExecution:** Ensures that pods cannot be scheduled on nodes with the specified labels (node-group in this case).

**NotIn:** Prevents scheduling on nodes with labels matching the specified values.

I upgraded the helm chart (first with a dry run) to test the changes:

```
$helm upgrade ping ./ --namespace default --dry-run --debug
```

No issues found I proceeded with the upgrade :

```
$helm upgrade ping ./ --namespace default
```

Release "ping" has been upgraded. Happy Helming!

NAME: ping

LAST DEPLOYED: Sat Jan 18 11:45:53 2025

NAMESPACE: default

STATUS: deployed

REVISION: 5

NOTES:

1. Get the application URL by running these commands:

```

export POD_NAME=$(kubectl get pods --namespace default -l
"app.kubernetes.io/name=ping,app.kubernetes.io/instance=ping" -o jsonpath="{.items[0].metadata.name}")
export CONTAINER_PORT=$(kubectl get pod --namespace default $POD_NAME -o
jsonpath="{.spec.containers[0].ports[0].containerPort}")
echo "Visit http://127.0.0.1:8080 to use your application"
kubectl --namespace default port-forward $POD_NAME 8080:$CONTAINER_PORT

```

```

$kubectl get pods -n default
NAME                READY STATUS RESTARTS AGE
ping-57d5b87fd6-2zfzb 1/1   Running 0      107s

```

```

$kubectl describe pods ping-57d5b87fd6-2zfzb -n default
See output in [5]

```

```

$kubectl get pod ping-57d5b87fd6-2zfzb -n default -o yaml | grep -A 5 nodeAffinity
nodeAffinity:
  requiredDuringSchedulingIgnoredDuringExecution:
    nodeSelectorTerms:
    - matchExpressions:
      - key: node-group
        operator: NotIn

```

The pod will not be scheduled on any node with `node-group` labels matching the forbidden values specified in `values.yaml`.

Let's test it further :

I temporarily labeled all the nodes with a forbidden value.

```

kubectl label node gke-cluster-1-default-pool-1b47b03f-0k57 node-group=forbidden-group-1
kubectl label node gke-cluster-1-default-pool-1b47b03f-1k33 node-group=forbidden-group-1
kubectl label no gke-cluster-1-default-pool-1b47b03f-vn8v node-group=forbidden-group-2

```

```

node/gke-cluster-1-default-pool-1b47b03f-0k57 labeled
...
...

```

Checking the current pod:

```

kubectl get pod ping-57d5b87fd6-7nm4m -n default -o wide
NAME                READY STATUS RESTARTS AGE IP      NODE                                NOMINATED
NODE READINESS GATES
ping-57d5b87fd6-7nm4m 1/1   Running 0      17m  10.100.1.5  gke-cluster-1-default-pool-1b47b03f-1k33  <none>
<none>

```

I forced rescheduling by deleting and recreating the pod:

```

kubectl delete pod ping-57d5b87fd6-7nm4m -n default
pod "ping-57d5b87fd6-7nm4m" deleted

```

Checking the current pod I can see :

```

kubectl get pod ping-57d5b87fd6-m7s7n -n default -o wide
NAME                READY STATUS RESTARTS AGE IP      NODE                                NOMINATED NODE READINESS
GATES
ping-57d5b87fd6-m7s7n 0/1   Pending 0      2m16s <none> <none> <none> <none>

```

It's on pending status as expected, as pod didn't match the node affinity rules.

```

$kubectl describe pod ping-57d5b87fd6-m7s7n -n default

```

```

Warning FailedScheduling 3m16s default-scheduler 0/3 nodes are available: 3 node(s) didn't match Pod's node
affinity/selector. precondition: 0/3 nodes are available: 3 Preemption is not helpful for scheduling.

```

Full output here [6].

Test finished, affinity is working I will remove the node labels:

```
$kubectl label node gke-cluster-1-default-pool-1b47b03f-0k57 node-group-  
node/gke-cluster-1-default-pool-1b47b03f-0k57 unlabeled  
$kubectl label node gke-cluster-1-default-pool-1b47b03f-1k33 node-group-  
$kubectl label no gke-cluster-1-default-pool-1b47b03f-vn8v node-group-
```

## Challenge 1.2:

Ensure that a pod will not be scheduled on a node that already has a pod of the same type.

To ensure that a pod is not scheduled on a node that already has a pod of the same type, I can use a **pod anti-affinity** rule in the Helm chart.

I will add a podAntiAffinity rule under the affinity section of the deployment.yaml.

This ensures that pods with the same labels do not get scheduled on the same node.

Adding to deployment :

```
podAntiAffinity:  
  requiredDuringSchedulingIgnoredDuringExecution:  
  - labelSelector:  
      matchLabels:  
        app.kubernetes.io/name: {{ .Chart.Name }}  
    topologyKey: "kubernetes.io/hostname"
```

- **podAntiAffinity:** Specifies rules to avoid scheduling pods of the same type on the same node.
- **requiredDuringSchedulingIgnoredDuringExecution:** Ensures strict enforcement during scheduling. Pods will not schedule on nodes that already have matching pods.
- **labelSelector:** Matches pods with specific labels. In this case, it ensures the rule applies to pods with the app.kubernetes.io/name label matching the chart's name.
- **topologyKey:** Defines the scope for the rule. "kubernetes.io/hostname" enforces this rule on a per-node basis.

Let's render the yaml :

```
$helm template ./ --namespace default  
$helm upgrade --install ping ./ --namespace default
```

Let's increase the replica count in values.yaml

```
replicaCount: 4  
helm upgrade --install ping ./ --namespace default
```

Checking the labels :

```
kubectl get pods -n default -o yaml | grep -A 10 "labels"
```

Label “ping” is correct.

Checking the pods :

```
kubectl get pods -o wide -n default
NAME                READY   STATUS    RESTARTS   AGE   IP            NODE                                     NOMINATED
NODE READINESS GATES
ping-58b97c5bb4-2jts 0/1     Pending   0          76s   <none>        <none>                                <none>
<none>
ping-58b97c5bb4-8t4d 1/1     Running   0          76s   10.100.1.12   gke-cluster-1-default-pool-1b47b03f-1k33
<none> <none>
ping-58b97c5bb4-x6xhp 1/1     Running   0          76s   10.100.0.8    gke-cluster-1-default-pool-1b47b03f-vn8v <none>
<none>
ping-58b97c5bb4-xt6dx 1/1     Running   0          76s   10.100.2.18   gke-cluster-1-default-pool-1b47b03f-0k57 <none>
<none>
```

```
$kubectl describe pod ping-58b97c5bb4-2jts -n default
```

Events:

Type	Reason	Age	From	Message
Warning	FailedScheduling	2m28s	default-scheduler	0/3 nodes are available: 3 node(s) didn't match pod anti-affinity rules. preemption: 0/3 nodes are available: 3 No preemption victims found for incoming pod.
Warning	FailedScheduling	2m27s	default-scheduler	0/3 nodes are available: 3 node(s) didn't match pod anti-affinity rules. preemption: 0/3 nodes are available: 3 No preemption victims found for incoming pod.
Normal	NotTriggerScaleUp	2m29s	cluster-autoscaler	pod didn't trigger scale-up:

Challenge 1.2 completed. Pod anti-affinity rules are working.

---

## Challenge 1.3

Ensure that the Pods are deployed across different availability zones.

To ensure pods are deployed across different availability zones, I will modify the deployment.yaml to include a **podAntiAffinity** rule based on the zone topology. Kubernetes uses the topology.kubernetes.io/zone label to represent the availability zone of each node.

Changes to deployment.yaml to include a podAntiAffinity rule targeting zones :

```
#Ensure pods are spread across nodes (anti-affinity based on hostname)
podAntiAffinity:
  requiredDuringSchedulingIgnoredDuringExecution:
    - labelSelector:
        matchLabels:
          app.kubernetes.io/name: {{ .Chart.Name }}
        topologyKey: "kubernetes.io/hostname" # Avoid scheduling multiple pods of the same type on the same node
    #Avoid scheduling in the same zone
    - label selector:
        matchLabels:
          app.kubernetes.io/name: {{ .Chart.Name }}
        topologyKey: "topology.kubernetes.io/zone"
```

I currently have a zonal cluster, I will create a new cluster to be able to test it.

In a zonal cluster it works as expected of 4 replicas only one is scheduled:

NAME	READY	STATUS	RESTARTS	AGE
ping-6cb9b69695-6tdvb	0/1	Pending	0	83m
ping-6cb9b69695-7bw8f	1/1	Running	0	83m
ping-6cb9b69695-7p9gc	0/1	Pending	0	83m
ping-6cb9b69695-h88lv	0/1	Pending	0	83m

Trying it in a regional cluster [7]:

```
$helm install ping ./ --namespace default
```

NAME	READY	STATUS	RESTARTS	AGE
ping-6cb9b69695-7nwm4	1/1	Running	0	6s
ping-6cb9b69695-7x6j8	0/1	Pending	0	6s
ping-6cb9b69695-w274f	1/1	Running	0	6s
ping-6cb9b69695-w6dxm	1/1	Running	0	6s

Expected result, one pod in each zone, the solution is working.

---

## Challenge 1.4

Ensure that another random service is up before applying the manifests.

I will create a pre-install hook to check for the random service.

templates/hooks/check-random-service.yaml that will ensure the random service is up before proceeding [8].

The helm chart structure will be like :

```
ping/
├── Chart.yaml
├── values.yaml
├── templates/
│   ├── deployment.yaml
│   ├── service.yaml
│   └── hooks/
│       └── check-random-service.yaml
```

I initially thought in creating the random service as a dependency of the ping helm chart and loading it in the Chart.yaml. However, after testing it as a dependency, the hook was created before the random-service creation. I ended up with a simple workaround, I created 2 separate Helm Releases, because I wanted to be completely sure that random service was created before the hook and the ping deployment.

1. Deployed random-service first:

```
$helm install random-service ./charts/random-service --namespace default --debug
```

This release contains the random service.

2. After the random service I deployed “ping”:

This release contains the ping workload with a prehook that will check if the random workload is running (pre-install) and continue the ping deployment if it is running.

```
helm install ping ./ --namespace default --debug
```

This ensures random-service is fully deployed before the ping chart.

The random-service helm code is :

Deployment -> ping/charts/random-service/templates/deployment.yaml [9]

Service -> ping/charts/random-service/templates/service.yaml [10]

The hook code is :

ping/templates/hooks/check-random-service.yaml [8]

I tested it in the already created GKE cluster, the logic between the hook works and only if random-service is up , the ping workload is deployed.

---

## Challenge 2

To automate the described process of copying Helm charts from the reference GCP Artifact Registry (reference.gcr.io) to another GCP Artifact Registry (instance.gcr.io), I'll provide an automation solution using **Terraform** module.

```
terraform-project/
├─ main.tf           # Main Terraform configuration file
├─ variables.tf
├─ terraform.tfvars
├─ modules/
│   └─ helm_chart_copy/
│       ├── main.tf   # Module implementation
│       ├── variables.tf # Variables specific to the module
│       └─ outputs.tf  # Outputs specific to the module
```

**Root main.tf** (Main Terraform configuration file where the module is invoked.)

```
terraform {
  required_providers {
    google = {
      source = "hashicorp/google"
      version = "~> 4.0"
    }
  }

  required_version = ">= 1.3.0"
}

provider "google" {
  credentials = file(var.service_account_key)
  project     = var.project_id
  region      = var.region
}

module "copy_helm_charts" {
  source          = "./modules/helm_chart_copy"
  source_registry = var.source_registry
  target_registry = var.target_registry
  helm_chart_list = var.helm_chart_list
  sa_email        = var.sa_email
}
```

**Root Variables.tf** (all configurable inputs in `main.tf` are parameterized, promoting reusability and flexibility)

```
variable "project_id" {
  description = "The ID of the Google Cloud project"
  type        = string
}

variable "region" {
  description = "The region for the Google Cloud resources"
  type        = string
}

variable "service_account_key" {
  description = "Path to the service account key JSON file"
  type        = string
}
```



```

variable "helm_chart_list" {
  description = "List of Helm charts to copy"
  type       = list(string)
}

variable "source_registry" {
  description = "The source registry for Helm charts"
  type       = string
  default    = "https://reference.gcr.io/helm"
}

variable "target_registry" {
  description = "The target registry for Helm charts"
  type       = string
}

variable "sa_email" {
  description = "The email of the service account to use for authentication"
  type       = string
}

```

## Root Terraform.tfvars example file

```

project_id    = "your-project-id"
region       = "us-central1"
service_account_key = "/path/to/your/service-account-key.json"
helm_chart_list = ["chart1", "chart2"]
source_registry = "https://reference.gcr.io/helm"
target_registry = "instance.gcr.io"
sa_email      = "your-service-account-email"

```

## Module: modules/helm\_chart\_copy/main.tf

#The logic for copying Helm charts resides here.

```

variable "source_registry" {}
variable "target_registry" {}
variable "helm_chart_list" {
  type = list(string)
}
variable "sa_email" {}

resource "null_resource" "copy_helm_charts" {
  provisioner "local-exec" {
    command = <<EOT
    gcloud auth activate-service-account ${var.sa_email} --key-file=/path/to/key.json

    # Add source Helm repository
    helm repo add source ${var.source_registry}
    helm repo update

    for chart in ${join(" ", var.helm_chart_list)}; do
      # Pull the Helm chart from the source repository
      helm pull source/$chart --untar

      # Push the chart to the target Artifact Registry
      helm push $chart.tar.gz oci://${var.target_registry}
    done
  EOT
}

```

### Module: modules/helm\_chart\_copy/variables.tf

The modules/helm\_chart\_copy/variables.tf file defines variables specifically for the **module**, making the module independent and reusable.

```
variable "source_registry" {  
  description = "Source Helm repository URL"  
  type        = string  
}  
variable "target_registry" {  
  description = "Target GCP Artifact Registry URL"  
  type        = string  
}  
variable "helm_chart_list" {  
  description = "List of Helm charts to copy"  
  type        = list(string)  
}  
variable "sa_email" {  
  description = "Service account email for permissions"  
  type        = string  
}
```

### Module: modules/helm\_chart\_copy/outputs.tf (Optional)

Useful information about the process using an outputs file.

```
output "copied_helm_charts" {  
  description = "List of Helm charts copied"  
  value       = var.helm_chart_list  
}
```

### Module Code Explanation:

**helm pull:** This downloads the chart directly from the source Helm repository. It avoids the need for using fetch for OCI repositories.

**Helm Repository Management:** Adds and updates the Helm repository before pulling charts.

**OCI Push:** Assumes the target registry is OCI-compliant and pushes the charts accordingly.

### How to use the module:

Replace placeholders with the correct values (source\_registry, target\_registry, etc.).

Ensure the service account (sa) has necessary roles in GCP:

roles/artifactregistry.reader for the source repository.

roles/artifactregistry.writer for the target repository.

Save this as a module and call it from the Terraform configuration.

## Challenge 3

Create a Github workflow to allow installing helm chart from Challenge #1 using module from Challenge #2, into a GKE cluster (considering a preexisting resource group and cluster name).

I create the file structure:

```
.github/  
└─ workflows/  
    └─ install-helm-chart.yaml
```

`.github/workflows/install-helm-chart.yaml`.

The install-helm-chart.yaml workflow yaml content:

```
name: Deploy Helm Chart to GKE

on:
  workflow_dispatch: # Allows manual triggering of the workflow
  push:
    branches:
      - main # Trigger on push to the main branch

jobs:
  deploy-helm:
    runs-on: ubuntu-latest

    env:
      GOOGLE_CREDENTIALS: ${ secrets.GOOGLE_CREDENTIALS } # GCP Service Account Key
      PROJECT_ID: ${ secrets.GCP_PROJECT_ID } # GCP Project ID
      REGION: ${ secrets.GCP_REGION } # GCP Region
      CLUSTER_NAME: ${ secrets.GKE_CLUSTER_NAME } # GKE Cluster Name

    steps:
      - name: Checkout Code
        uses: actions/checkout@v3

      - name: Set up Google Cloud SDK
        uses: google-github-actions/setup-gcloud@v1
        with:
          service_account_key: ${ secrets.GOOGLE_CREDENTIALS }
          project_id: ${ secrets.GCP_PROJECT_ID }

      - name: Authenticate with GKE
        run: |
          set -e
          gcloud container clusters get-credentials $CLUSTER_NAME --region $REGION --project $PROJECT_ID

      - name: Set up Terraform
        uses: hashicorp/setup-terraform@v2
        with:
          terraform_wrapper: true

      - name: Initialize Terraform
        run: terraform init

      - name: Plan Terraform Changes
        run: |
          set -e
          terraform plan \
            -var="project_id=${ secrets.GCP_PROJECT_ID }" \
            -var="region=${ secrets.GCP_REGION }" \
            -var="cluster_name=${ secrets.GKE_CLUSTER_NAME }"
```

```

- name: Apply Terraform Changes
run: |
  set -e
  terraform apply -auto-approve \
    -var="project_id=${{ secrets.GCP_PROJECT_ID }}" \
    -var="region=${{ secrets.GCP_REGION }}" \
    -var="cluster_name=${{ secrets.GKE_CLUSTER_NAME }}"

- name: Verify Deployment
run: |
  kubectl get pods -n default

```

## Secrets required in the GitHub repository:

1. **GOOGLE\_CREDENTIALS:**  
JSON key file for a GCP service account with the following permissions:  
Kubernetes Engine Admin  
Artifact Registry Reader  
Service Account User
2. **GCP\_PROJECT\_ID:**  
The GCP project ID where the cluster resides.
3. **GCP\_REGION:**  
The region of the GKE cluster.
4. **GKE\_CLUSTER\_NAME:**  
The name of the existing GKE cluster.

To test the GitHub action:

```

$git add .github/workflows/install-helm-chart.yaml
$git commit -m "Add Helm chart deployment workflow"
$git push origin main

```

Select the workflow and click **Run Workflow**.

## Final Notes

Changed the ping chart version to 0.1.1.

I created the helm package:

```
$ helm package.
```

```
ping-0.1.1.tgz
```

Included also the helm release for the random-service as random-service-1.0.0.tgz.  
(you can find it also inside ping release in charts folder)

```
[1] $gcloud beta container --project "instant-jetty-435813-p4" clusters create "cluster-1" --
zone "us-central1-a" --tier "standard" --no-enable-basic-auth --cluster-version "1.30.8-
gke.1051000" --release-channel "regular" --machine-type "e2-medium" --image-type
"COS_CONTAINERD" --disk-type "pd-balanced" --disk-size "100" --metadata disable-legacy-
endpoints=true --scopes
"https://www.googleapis.com/auth/devstorage.read_only","https://www.googleapis.com/aut
h/logging.write","https://www.googleapis.com/auth/monitoring","https://www.googleapis.co
m/auth/servicecontrol","https://www.googleapis.com/auth/service.management.readonly","
https://www.googleapis.com/auth/trace.append" --num-nodes "3" --
logging=SYSTEM,WORKLOAD --
monitoring=SYSTEM,STORAGE,POD,DEPLOYMENT,STATEFULSET,DAEMONSET,HPA,CADVISOR,
KUBELET --enable-ip-alias --network "projects/instant-jetty-435813-
p4/global/networks/default" --subnetwork "projects/instant-jetty-435813-p4/regions/us-
central1/subnetworks/default" --no-enable-intra-node-visibility --default-max-pods-per-node
"110" --enable-ip-access --security-posture=standard --workload-vulnerability-
scanning=disabled --no-enable-master-authorized-networks --no-enable-google-cloud-access --
addons HorizontalPodAutoscaling,HttpLoadBalancing,GcePersistentDiskCsiDriver --enable-
autoupgrade --enable-autorepair --max-surge-upgrade 1 --max-unavailable-upgrade 0 --
binauthz-evaluation-mode=DISABLED --enable-managed-prometheus --enable-shielded-nodes
--node-locations "us-central1-a"
[2] kubectl cluster-info
```

Kubernetes control plane is running at https://104.154.163.47

GLBCDefaultBackend is running at https://104.154.163.47/api/v1/namespaces/kube-
system/services/default-http-backend:http/proxy

KubeDNS is running at https://104.154.163.47/api/v1/namespaces/kube-
system/services/kube-dns:dns/proxy

Metrics-server is running at https://104.154.163.47/api/v1/namespaces/kube-
system/services/https:metrics-server:/proxy

```
[3] {
  "host": {
    "hostname": "127.0.0.1",
    "ip": "::ffff:127.0.0.1",
    "ips": []
  },
  "http": {
    "method": "GET",
    "baseUrl": "",
    "originalUrl": "/?authuser=0",
```

```
"protocol": "http"
},
"request": {
  "params": {
    "0": "/"
  },
  "query": {
    "authuser": "0"
  },
  "cookies": {

  },
  "body": {

  },
  "headers": {
    "host": "127.0.0.1:8080",
    "user-agent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/131.0.0.0 Safari/537.36",
    "accept": "text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.7",
    "accept-encoding": "gzip, deflate, br, zstd",
    "accept-language": "es-ES,es;q=0.9,ca;q=0.8",
    "referrer": "https://shell.cloud.google.com/",
    "sec-ch-ua": "\"Google Chrome\";v=\"131\", \"Chromium\";v=\"131\", \"Not_A Brand\";v=\"24\"",
    "sec-ch-ua-mobile": "?0",
    "sec-ch-ua-platform": "\"Windows\"",
    "sec-fetch-dest": "document",
    "sec-fetch-mode": "navigate",
    "sec-fetch-site": "cross-site",
```

```
"sec-user-ip": "10.0.0.32",

"upgrade-insecure-requests": "1",

"x-forwarded-host": "8080-cs-848137978281-default.cs-europe-west1-
onse.cloudshell.dev"

}

},

"environment": {

"PATH": "/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin",

"HOSTNAME": "ping-5989f667d5-624mc",

"NODE_VERSION": "16.16.0",

"YARN_VERSION": "1.22.19",

"KUBERNETES_PORT_443_TCP": "tcp://34.118.224.1:443",

"PING_SERVICE_PORT_HTTP": "80",

"PING_PORT": "tcp://34.118.229.251:80",

"PING_PORT_80_TCP_ADDR": "34.118.229.251",

"KUBERNETES_PORT": "tcp://34.118.224.1:443",

"KUBERNETES_PORT_443_TCP_ADDR": "34.118.224.1",

"PING_PORT_80_TCP_PROTO": "tcp",

"PING_PORT_80_TCP_PORT": "80",

"KUBERNETES_SERVICE_PORT": "443",

"KUBERNETES_SERVICE_PORT_HTTPS": "443",

"PING_SERVICE_HOST": "34.118.229.251",

"PING_SERVICE_PORT": "80",

"KUBERNETES_SERVICE_HOST": "34.118.224.1",

"KUBERNETES_PORT_443_TCP_PROTO": "tcp",

"KUBERNETES_PORT_443_TCP_PORT": "443",

"PING_PORT_80_TCP": "tcp://34.118.229.251:80",

"HOME": "/root"

}

}
```

[4]: <https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/>

[5] Name: ping-57d5b87fd6-2zfzb  
Namespace: default  
Priority: 0  
Service Account: ping  
Node: gke-cluster-1-default-pool-1b47b03f-1k33/10.128.0.34  
Start Time: Sat, 18 Jan 2025 11:45:55 +0000  
Labels: app.kubernetes.io/instance=ping  
app.kubernetes.io/name=ping  
pod-template-hash=57d5b87fd6  
Annotations: <none>  
Status: Running  
IP: 10.100.1.5  
IPs: 10.100.1.5  
Controlled By: ReplicaSet/ping-57d5b87fd6  
Containers:  
ping:  
Container ID: containerd://9e1d41df62e9183e353b77cd88eb1010b339ccb8c45c9aae8049640a6e29fc55  
Image: ealen/echo-server:0.7.0  
Image ID: docker.io/ealen/echo-server@sha256:086cba978dd74f4bbdbe91230c929d64b77201767fcec07d12697ff28abffbf9  
Port: 80/TCP  
Host Port: 0/TCP  
State: Running  
Started: Sat, 18 Jan 2025 11:45:56 +0000  
Ready: True  
Restart Count: 0  
Liveness: http-get http://:http/ delay=0s timeout=1s period=10s #success=1 #failure=3  
Readiness: http-get http://:http/ delay=0s timeout=1s period=10s #success=1 #failure=3  
Environment: <none>  
Mounts:  
/var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-q64tr (ro)  
Conditions:  
Type Status  
PodReadyToStartContainers True  
Initialized True  
Ready True  
ContainersReady True  
PodScheduled True  
Volumes:  
kube-api-access-q64tr:  
Type: Projected (a volume that contains injected data from multiple sources)  
TokenExpirationSeconds: 3607  
ConfigMapName: kube-root-ca.crt  
ConfigMapOptional: <nil>  
DownwardAPI: true  
QoS Class: BestEffort  
Node-Selectors: <none>  
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s  
node.kubernetes.io/unreachable:NoExecute op=Exists for 300s  
Events:  
Type Reason Age From Message  
----  
Normal Scheduled 3m default-scheduler Successfully assigned default/ping-57d5b87fd6-2zfzb to gke-cluster-1-default-pool-1b47b03f-1k33  
Normal Pulled 3m kubelet Container image "ealen/echo-server:0.7.0" already present on machine  
Normal Created 3m kubelet Created container ping  
Normal Started 3m kubelet Started container ping

[6]: Events:

Name: ping-57d5b87fd6-m7s7n  
Namespace: default  
Priority: 0  
Service Account: ping  
Node: <none>  
Labels: app.kubernetes.io/instance=ping  
app.kubernetes.io/name=ping  
pod-template-hash=57d5b87fd6  
Annotations: cloud.google.com/cluster\_autoscaler\_unhelpable\_since: 2025-01-18T12:09:42+0000  
cloud.google.com/cluster\_autoscaler\_unhelpable\_until: Inf  
Status: Pending  
IP:  
IPs: <none>  
Controlled By: ReplicaSet/ping-57d5b87fd6  
Containers:  
ping:



Image: ealen/echo-server:0.7.0  
 Port: 80/TCP  
 Host Port: 0/TCP  
 Liveness: http-get http://:http/ delay=0s timeout=1s period=10s #success=1 #failure=3  
 Readiness: http-get http://:http/ delay=0s timeout=1s period=10s #success=1 #failure=3  
 Environment: <none>  
 Mounts:  
 /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-4fp42 (ro)  
 Conditions:  
 Type Status  
 PodScheduled False  
 Volumes:  
 kube-api-access-4fp42:  
 Type: Projected (a volume that contains injected data from multiple sources)  
 TokenExpirationSeconds: 3607  
 ConfigMapName: kube-root-ca.crt  
 ConfigMapOptional: <nil>  
 DownwardAPI: true  
 QoS Class: BestEffort  
 Node-Selectors: <none>  
 Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s  
 node.kubernetes.io/unreachable:NoExecute op=Exists for 300s  
 Events:  

Type	Reason	Age	From	Message
Warning	FailedScheduling	3m16s	default-scheduler	0/3 nodes are available: 3 node(s) didn't match Pod's node affinity/selector. preemption: 0/3 nodes are available: 3 Preemption is not helpful for scheduling.
Warning	FailedScheduling	3m15s	default-scheduler	0/3 nodes are available: 3 node(s) didn't match Pod's node affinity/selector. preemption: 0/3 nodes are available: 3 Preemption is not helpful for scheduling.
Normal	NotTriggerScaleUp	3m17s	cluster-autoscaler	pod didn't trigger scale-up:

```
[7] gcloud beta container --project "instant-jetty-435813-p4" clusters create
"cluster-2" --region "us-central1" --tier "standard" --no-enable-basic-auth --
cluster-version "1.30.8-gke.1051000" --release-channel "regular" --machine-
type "e2-medium" --image-type "COS_CONTAINERD" --disk-type "pd-balanced" --
disk-size "100" --metadata disable-legacy-endpoints=true --scopes
"https://www.googleapis.com/auth/devstorage.read_only","https://www.googleapis
.com/auth/logging.write","https://www.googleapis.com/auth/monitoring","https:/
/www.googleapis.com/auth/servicecontrol","https://www.googleapis.com/auth/serv
ice.management.readonly","https://www.googleapis.com/auth/trace.append" --num-
nodes "1" --logging=SYSTEM,WORKLOAD --
monitoring=SYSTEM,STORAGE,POD,DEPLOYMENT,STATEFULSET,DAEMONSET,HPA,CADVISOR,KU
BELET --enable-ip-alias --network "projects/instant-jetty-435813-
p4/global/networks/default" --subnetwork "projects/instant-jetty-435813-
p4/regions/us-central1/subnetworks/default" --no-enable-intra-node-visibility
--default-max-pods-per-node "110" --enable-ip-access --security-
posture=standard --workload-vulnerability-scanning=disabled --no-enable-
master-authorized-networks --no-enable-google-cloud-access --addons
HorizontalPodAutoscaling,HttpLoadBalancing,GcePersistentDiskCsiDriver --
enable-autoupgrade --enable-autorepair --max-surge-upgrade 1 --max-
unavailable-upgrade 0 --binauthz-evaluation-mode=DISABLED --enable-managed-
prometheus --enable-shielded-nodes
```

```
[8]: check-random-service.yaml
apiVersion: batch/v1
kind: Job
metadata:
  name: {{ .Release.Name }}-check-random-service
  annotations:
    "helm.sh/hook": pre-install,pre-upgrade
```

```

    "helm.sh/hook-delete-policy": hook-succeeded
spec:
  template:
    spec:
      containers:
        - name: check-random-service
          image: busybox
          command:
            - /bin/sh
            - -c
            - |
              for i in {1..60}; do
                if wget -qO- random-service-random1.default.svc.cluster.local:80; then
                  echo "Random service is up!";
                  exit 0;
                fi;
                echo "Waiting for random service...";
                sleep 5;
              done;
              echo "Random service is not ready.";
              exit 1;
          restartPolicy: Never
      backoffLimit: 1

```

[9] random-service deployment.yaml

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ .Release.Name }}-random1
  labels:
    app.kubernetes.io/name: {{ .Chart.Name }}
    app.kubernetes.io/instance: {{ .Release.Name }}
spec:
  replicas: {{ .Values.replicaCount }}
  selector:
    matchLabels:
      app.kubernetes.io/name: {{ .Chart.Name }}
  template:
    metadata:
      labels:
        app.kubernetes.io/name: {{ .Chart.Name }}
    spec:
      containers:
        - name: random-service
          image: "us-docker.pkg.dev/google-samples/containers/gke/hello-app:2.0"

```

10: Random-service service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: {{ .Release.Name }}-random1
  labels:
    app.kubernetes.io/name: {{ .Chart.Name }}
    app.kubernetes.io/instance: {{ .Release.Name }}
spec:
  type: ClusterIP
  ports:
    - port: 80
      targetPort: 8080
      protocol: TCP
  selector:
    app.kubernetes.io/name: {{ .Chart.Name }}
```